

Course Database Constraints and Transact-SQL Implementation

Constraints

A constraint is a *restriction*

- Ensures data integrity

Constraint types:

Domain constraints


- On columns

Entity constraints

- On records as a whole

Referential integrity constraints

- On relationships *between* records



And constraints to realise complex business rules

- by database *programming*

Constraints – Example from Adventureworks

User defined
data type

```
CREATE TYPE Flag FROM bit NOT NULL
```

```
CREATE TABLE Employee (  
    BusinessEntityID int NOT NULL,  
    JobTitle nvarchar(50) NOT NULL,  
    BirthDate date NOT NULL,  
    MaritalStatus nchar(1) NOT NULL,  
    Gender nchar(1) NOT NULL,  
    HireDate date NOT NULL,  
    SalariedFlag Flag NOT NULL,  
    ModifiedDate datetime NOT NULL,  
    CONSTRAINT PK_Employee_BusinessEntityID PRIMARY KEY (BusinessEntityID)  
)
```

Domain

Entity

```
ALTER TABLE Employee WITH CHECK  
ADD CONSTRAINT CK_Employee_BirthDate  
CHECK ((BirthDate >= '1930-01-01'  
    AND  
    BirthDate <= dateadd(year,(-18), getdate()))  
)
```

Domain

Constraints – Example from Adventureworks

```
ALTER TABLE Employee  
    ADD CONSTRAINT FK_Employee_Person_BusinessEntityID  
FOREIGN KEY(BusinessEntityID) REFERENCES Person (BusinessEntityID)
```



RI

```
ALTER TABLE Employee WITH CHECK  
    ADD CONSTRAINT CK_Employee_Gender CHECK ((upper(Gender) = 'F' OR upper(Gender)= 'M'))
```



Domain

```
ALTER TABLE Employee  
    ADD CONSTRAINT DF_Employee_SalariedFlag DEFAULT (1) FOR SalariedFlag
```

```
ALTER TABLE Employee  
    ADD CONSTRAINT DF_Employee_ModifiedDat DEFAULT (getdate()) FOR ModifiedDate
```

Entity and Domain Constraints

Entity Constraint

- requires every row to have a unique value for some column or combination of columns
- PRIMARY KEY (Implicit NOT NULL)
- UNIQUE (allows *one* NULL value)

Domain constraint

- requires that a particular column (or set of columns) meets some criteria
- data type
- required data (NULL / NOT NULL)
- CHECK
- DEFAULT

DEFAULT

Specifies which value must be inserted in a column when the column is not listed in the INSERT-statement

```
CREATE TABLE test_defaults (  
    keycol          smallint,  
    process_id      smallint      DEFAULT @@SPID,  
    date_ins        datetime      DEFAULT GETDATE(),  
    mathcol         smallint      DEFAULT 10 * 2,  
    text1           char(3),  
    text2           char(3)      DEFAULT 'xyz'  
)
```

```
INSERT test_defaults (text1) VALUES ('qqq')
```

```
SELECT * FROM test_defaults
```

keycol	process_id	date_ins	mathcol	text1	text2
NULL	52	2016-02-03 14:06:55.417	20	qqq	xyz

Referential Integrity (RI) constraint

**Requires that a value in one column matches
the value in another column**

- ‘child’ table refers to the ‘parent’ table
- in a different table or in the same table
 - maybe in another database
 - maybe on another server

RI constraint

Be aware that RI constraints can place restrictions on *both* tables

- is it allowed to delete/update a parent record if that record is referenced from a child?
- if so, what happens with the child?
 - CASCADING DELETE/UPDATE?
 - or is the reference SET to NULL
 - or is the reference SET to a DEFAULT?
 - or is the DELETE/UPDATE RESTRICTED?

Cardinality constraints

1:1, 1:N, N:M, 1-5, ...

(Mandatory/Optional) NULL/NOT NULL

But how to enforce that a soccer team has no more than 11 players?

- Use a stored procedure or trigger....

Exercise

Workbook, Theme CONSTRAINTS

Exercises 1 and 2

Taxonomy of Constraints - 1

static constraints

- rules which every record *at any instance of time* must obey
- ex: Domain constraints
 - Gender of a Person must be 'M' or 'F'
 - Data types
 - *limits the possible values (numeric, character), but that is usually not sufficient!*

dynamic constraints

- rules which *restrict the transition* between two instantiations of the database
 - ex: It is not allowed to update the marital status of a Person from 'Not Married' to 'Divorced'

Taxonomy of Constraints - 2

declarative constraints

- constraints are ‘declared’ by means of *DDL*- or *DCL*-statements
- specify **WHAT** has to be done
 - CHECK, NULL/NOT NULL, FOREIGN KEY, PRIMARY KEY, UNIQUE
- DBMS decides **HOW**

procedural constraints

- developer writes *code* to enforce constraints
 - triggers, stored procedures, user defined functions
- developer decides **HOW**

Why procedural constraints?

use declarative constraints whenever possible!

- fast, correct

but there are many things that DDL-statements can't specify

- complex relationships
- state transitions
- vendor specific:
 - RI across databases
 - Cascading updates/deletes in some cases
 - ...

but coding is non standard

- MS SQL Server: Transact-SQL (T-SQL)

Transact-SQL

Database Server

Most servers provide SQL-89 level functionality

Most servers include some SQL-92 features

Quite a few servers offer proprietary versions of SQL3 stored procedures, triggers, and rules

SQL Extensions

Relational DBMS's often have built-in procedural extensions to standard SQL

- A procedural programming language
 - variables, if-statements, while-statements, ...
- Stored procedures
- Triggers
- Rules

Useful, but non-standard

Vendor-specific SQL is no longer pure declarative

- Declarative: specify WHAT has to be done
- Procedural: specify HOW it has to be done
- MS SQL Server, Sybase: TRANSACT-SQL (T-SQL)
- Oracle: PL/SQL

Transact-SQL

Native language of SQL Server

- Extension of SQL with programming constructs

T-SQL is simple and limited

- Local variables
- Control of flow: IF, CASE, WHILE
- Operators, special functions
- **No** user interface programming constructs
- **No** file I/O

Program execution within the database engine

Consult *Books Online*!

Transact-SQL Syntax Conventions

Convention	Used for
UPPERCASE	Transact-SQL keywords.
<i>italic</i>	User-supplied parameters of Transact-SQL syntax.
bold	Database names, table names, column names, index names, stored procedures, utilities, data type names, and text that must be typed exactly as shown.
<u>underline</u>	Indicates the default value applied when the clause that contains the underlined value is omitted from the statement.
(vertical bar)	Separates syntax items enclosed in brackets or braces. You can use only one of the items.
[] (brackets)	Optional syntax items. Do not type the brackets.
{ } (braces)	Required syntax items. Do not type the braces.
[,... <i>n</i>]	Indicates the preceding item can be repeated <i>n</i> number of times. The occurrences are separated by commas.
[... <i>n</i>]	Indicates the preceding item can be repeated <i>n</i> number of times. The occurrences are separated by blanks.
[;]	Optional Transact-SQL statement terminator. Do not type the brackets.
<label> ::=	The name for a block of syntax. This convention is used to group and label sections of lengthy syntax or a unit of syntax that can be used in more than one location within a statement. Each location in which the block of syntax can be used is indicated with the label enclosed in chevrons: <label>.

T-SQL CREATE TABLE

(simplified)

CREATE TABLE

```
[ database_name.[schema name] . | schema name.] table_name  
(  
    { <column_definition> | computed_column_definition> }  
    [ <table_constraint>] [...n]  
    [;]  
)
```

```
<column_definition> ::= column_name <data_type>  
    [NULL | NOT NULL]  
    [  
        [ CONSTRAINT constraint_name] DEFAULT  
        constant_expression ]  
    | [ IDENTITY [(seed, increment) ] ]  
    ]  
    [ <column_constraint>] [ ...n]
```

T-SQL CREATE TABLE

(simplified)

```
<table_constraint> ::= [CONSTRAINT constraint_name]  
{  
    [ { PRIMARY KEY | UNIQUE }  
    [CLUSTERED | NONCLUSTERED]  
    { ( column[,...n] ) }  
]  
  
    | FOREIGN KEY  
      [ (column[,...n])]  
      REFERENCES ref_table [ (ref_column[,...n])] ]  
  
    | CHECK ( search_conditions )  
}
```

Note: Multicolumn key constraints are created as table constraints.

Multi-part names

Object reference format	Description
<i>server . database . schema . object</i>	Four-part name.
<i>server . database .. object</i>	Schema name is omitted.
<i>server .. schema . object</i>	Database name is omitted.
<i>server ... object</i>	Database and schema name are omitted.
<i>database . schema . object</i>	Server name is omitted.
<i>database .. object</i>	Server and schema name are omitted.
<i>schema . object</i>	Server and database name are omitted.
<i>object</i>	Server, database, and schema name are omitted.

For *Schema* study BOL, Transact-SQL Syntax Conventions,
Multi-part Names

Script, example

```
USE SomeDatabase
```

```
DECLARE @Ident INT -- declaration of variable
```

```
INSERT INTO Orders (CustomerID, OrderDate)  
VALUES (25, DATEADD(DAY,-1,GETDATE())) -- system functions
```

```
SELECT @Ident = @@IDENTITY /* variable assignment  
@@: system functions */
```

```
INSERT INTO Details (OrderID, ProductID, UnitPrice, Quantity)  
VALUES (@Ident, 1, 50, 25) -- retrieval of variable
```

```
SELECT 'The OrderID of the inserted row is ' + -- operator +  
CAST(@Ident AS varchar(8))
```

Variables

Declaration:

```
DECLARE @variable_name datatype
```

Assignment of a value:

```
SELECT @variable = expression  
[FROM ... [WHERE ...]]
```

Or

```
SET @variable = expression
```

Retrieval of a value:

```
SELECT @variable
```

For use in a batch, stored procedure or trigger or udf

- A variable exists only for the life of the batch

Arghh... Exercise

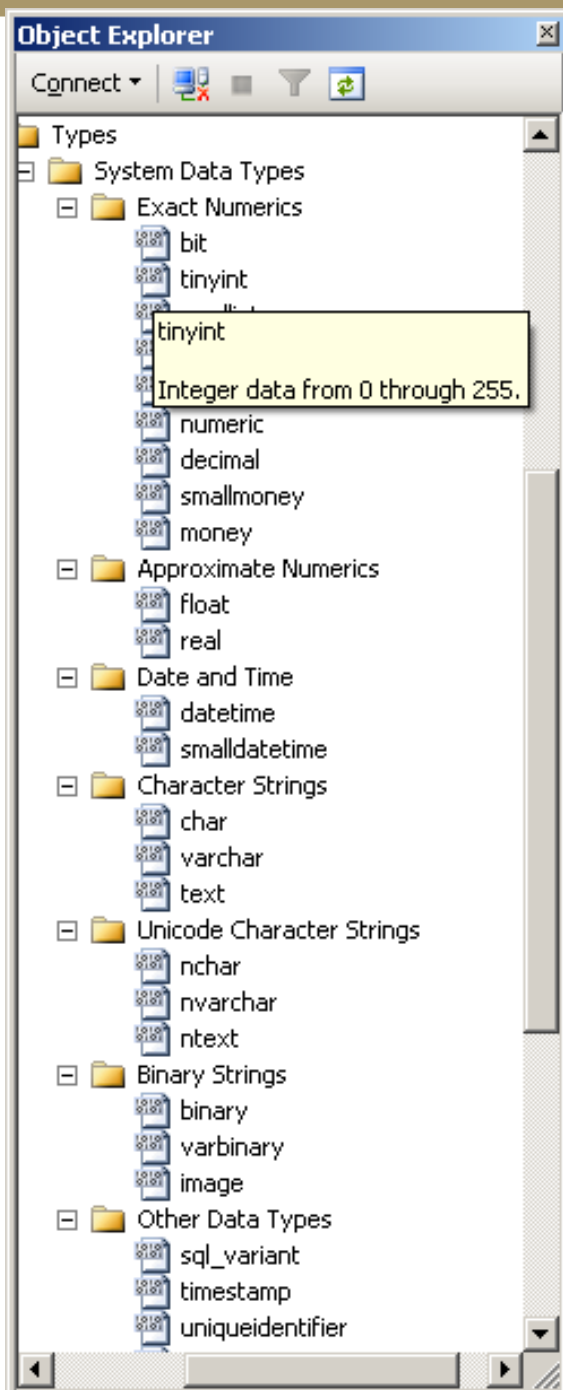
-- execute

```
DECLARE @composerName varchar(20)
```

```
SELECT @composerName = name  
FROM Composer
```

```
SELECT @composerName
```





T-SQL Data types

Study:

- Conversion table
- BOL
- Object Explorer

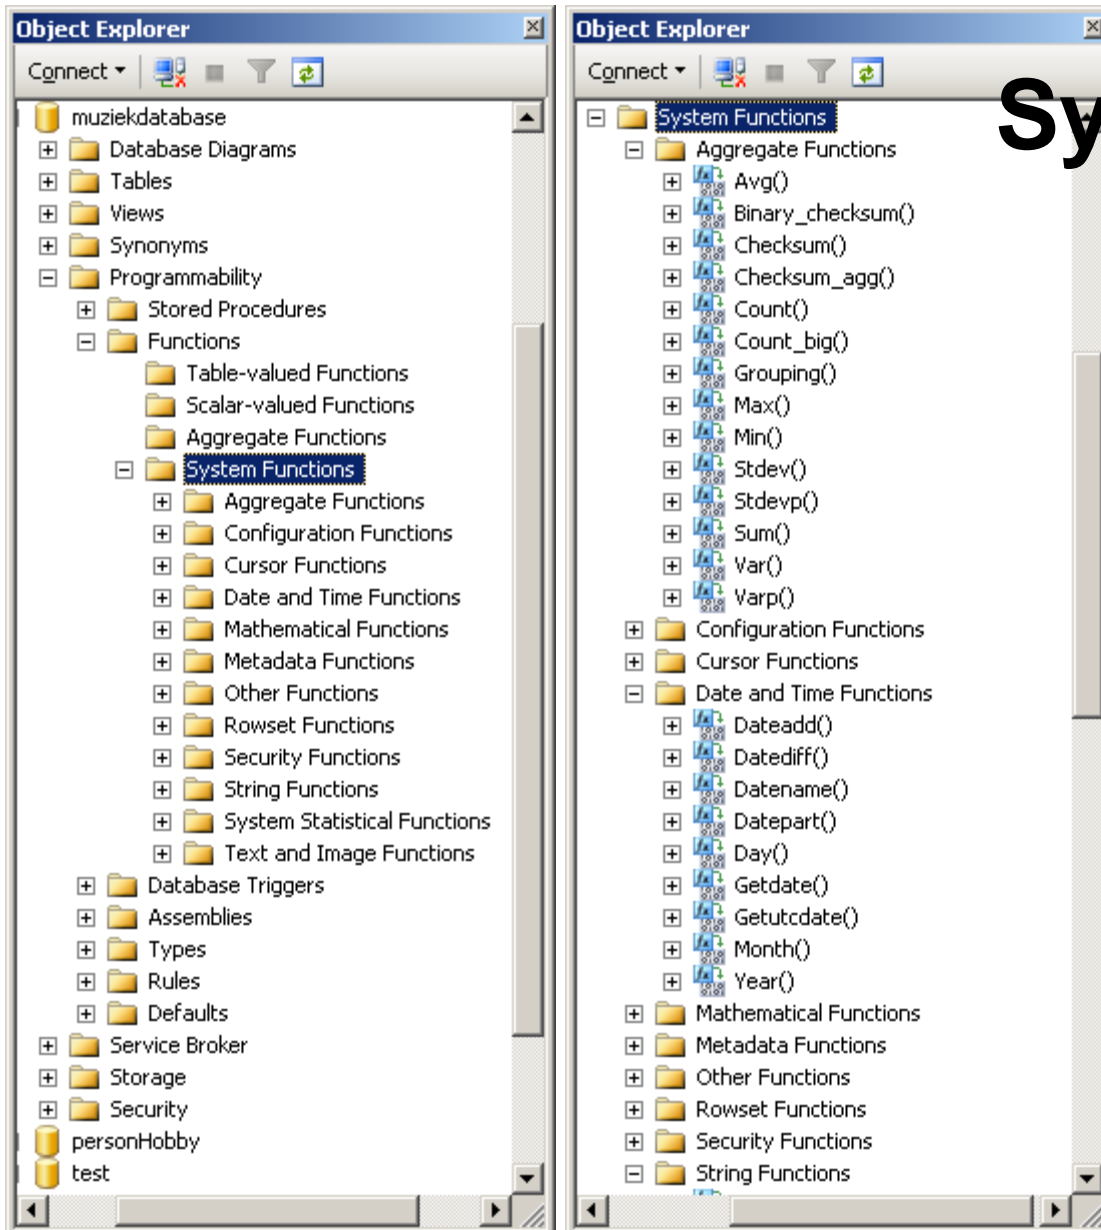
Identity Property

Creates an *identity column* in a table

Syntax: **IDENTITY** [(*seed* , *increment*)]

```
CREATE TABLE new_employees (  
    emp_id    INT PRIMARY KEY    IDENTITY(1,1),  
    fname     VARCHAR(20),  
    middle    CHAR(1),  
    lname     VARCHAR(30)  
)
```

```
INSERT new_employees (fname, middle, lname)  
VALUES ('Karin', 'F', 'Josephs')
```



System Functions

T-SQL Functions, Operators

Usage:

```
SELECT GETDATE()  
SELECT COS(PI())  
SELECT @my_pi = ROUND(PI(), 1)
```

Operators: +, -, *, /, %, =, >, <, >=, <=, <>, ...

Study Transact-SQL Help

Batch

A batch is a set of SQL-statements that is sent **at one time** to the server

The batch is by the server:

1. Parsed as a whole

- The server checks the syntax of each statement.

2. Optimised and compiled as a whole

- The server specifies the most efficient method to execute the queries

3. Executed statement by statement

Batch Example 1

A batch of statements of which one statement has a syntax error will not be executed

```
CREATE TABLE test  
  (col1 INT PRIMARY KEY,  
   col2 INT NOT NULL)
```

one
batch {
 INSERT test (col1, col2) VALUES (6,10)
 INSERT test (col1, col2) VLUES (6,51)
 UPDATE test SET col2 = 4
 SELECT * FROM test

Msg 102, Level 15, State 1, Line 2
Incorrect syntax near 'VLUES'.

Batch Example 2

If the server decides to execute the batch, it is of course possible that run-time errors occur

correct batch {
 INSERT test (col1, col2) VALUES (6,10)
 INSERT test (col1, col2) VALUES (6,51)
 UPDATE test SET col2 = 4
 SELECT * FROM test

(1 row(s) affected)

Msg 2627, Level 14, State 1, Line 2
Violation of PRIMARY KEY constraint 'PK__test__0BC6C43E'.
Cannot insert duplicate key in object 'dbo.test'.
The statement has been terminated.

(1 row(s) affected)

col1	col2
6	4

(1 row(s) affected)

T-SQL Local Variables

```
CREATE TABLE test2 (col INT NOT NULL)
```

```
INSERT test2 (col) VALUES (1)
```

```
DECLARE @maxcol INT
```

```
SELECT @maxcol = MAX(col) FROM test2
```

```
INSERT INTO test2 (col) VALUES (@maxcol + 1)
```

} execute as
one batch

After execution of this batch:

- the variable @maxcol is no longer defined

GO is NOT an SQL-stmt

If you want to send *two* batches to the server, you will have to do so one by one

Another way is to use a separator that instructs the *CLIENT* to send two batches

- A lot of front-end tools use the keyword **go**

T-SQL Conditional Execution

```
IF Boolean_expression  
    {Transact-sql_statement | statement_block}  
[ELSE [Boolean_expression]  
    {Transact-sql_statement | statement_block}]
```

```
statement_block:  
    BEGIN  
        Transact-SQL_statements  
    END
```

T-SQL Conditional Execution

```
CREATE TABLE product
( prod_nr      int          NOT NULL
  CONSTRAINT pk_product PRIMARY KEY (prod_nr),
  name         varchar(30)  NOT NULL,
  price        money        NOT NULL,
  type         varchar(30)  NOT NULL
)
```

```
INSERT product (prod_nr, name, price, type)
VALUES (1, 'tv', 500, 'electronics')
```

```
INSERT product (prod_nr, name, price, type)
VALUES (2, 'radio', 100, 'electronics')
```

```
INSERT product (prod_nr, name, price, type)
VALUES (3, 'ball', 100, 'sport')
```

Conditional Execution

Exercise 1

```
IF EXISTS (SELECT 1 FROM product WHERE name = 'radio')
  BEGIN
    PRINT 'There are radio''s in stock'
  END
ELSE
  BEGIN
    PRINT 'No radio''s in stock'
  END
```

Conditional Execution

Exercise 2

```
DECLARE @number INT
SET @number = 100
IF (SELECT AVG(price)
     FROM product WHERE type = 'electronics') > @number
BEGIN
    PRINT 'The average price of electronic products is greater than '
        + CONVERT(varchar(10), @number)
END
ELSE
BEGIN
    RAISERROR ('The average price of electronic products
              is less than %d.', 16, 1, @number)
END
```

T-SQL Repeated Execution

```
WHILE Boolean_expression  
    {Transact-sql_statement | statement_block}
```

```
WHILE (SELECT AVG(price) FROM product) < 500  
    BEGIN  
        UPDATE product  
        SET price = price * 1.05  
    END
```

Exercise

Modulo11-check:

972428577 is a valid bankaccountnumber, because:

$$(9*9 + 8*7 + 7*2 + 6*4 + 5*2 + 4*8 + 3*5 + 2*7 + 1*7) \% 11 = 0$$

Make a script that checks this for a fixed accountnumber

CASE, with *input expression*

```
USE AdventureWorks  
GO
```

```
SELECT TOP 10 SalesOrderID%10 AS 'OrderLastDigit',  
    ProductID%10 AS 'ProductLastDigit',  
    'How Close?' =  
        CASE SalesOrderID%10    -- input expression  
            WHEN ProductID%10    THEN 'Exact Match!'  
            WHEN ProductID%10 - 1 THEN 'Within 1'  
            WHEN ProductID%10 + 1 THEN 'Within 1'  
            ELSE 'More Than One Apart'  
        END  
FROM Sales.SalesOrderDetail  
ORDER BY SalesOrderID DESC
```


A searched CASE, without *input expression*

```
SELECT TOP 10 SalesOrderID%10 AS 'OrderLastDigit',  
ProductID%10 AS 'ProductLastDigit',  
'How Close?' =  
CASE  
WHEN (SalesOrderID%10) < 3 THEN 'Ends with less than three'  
WHEN ProductID = 6 THEN 'ProductID is 6'  
WHEN ABS(SalesOrderID%10 - ProductID) <= 1 THEN 'Within 1'  
ELSE 'More Than One Apart'  
END  
FROM Sales.SalesOrderDetail  
ORDER BY SalesOrderID DESC
```

Evaluates to
a Boolean

▶ H A N

Exercise

Use a searched CASE in a script that gives the name of the current day translated in Finnish.